

# MODULE 5

## Automating App Deployment

# Module Objectives

By the end of this module, you should understand:

- Interfaces into providing initialization of “cloud” instances
  - Cloud-Init
  - Puppet/Chef in a master-less environment
  - Use via CLI/Horizon
- Alternate Ways of initializing/managing instances
  - HEAT (and Ceilometer for triggering events)

# App Deployment - Then vs. Now

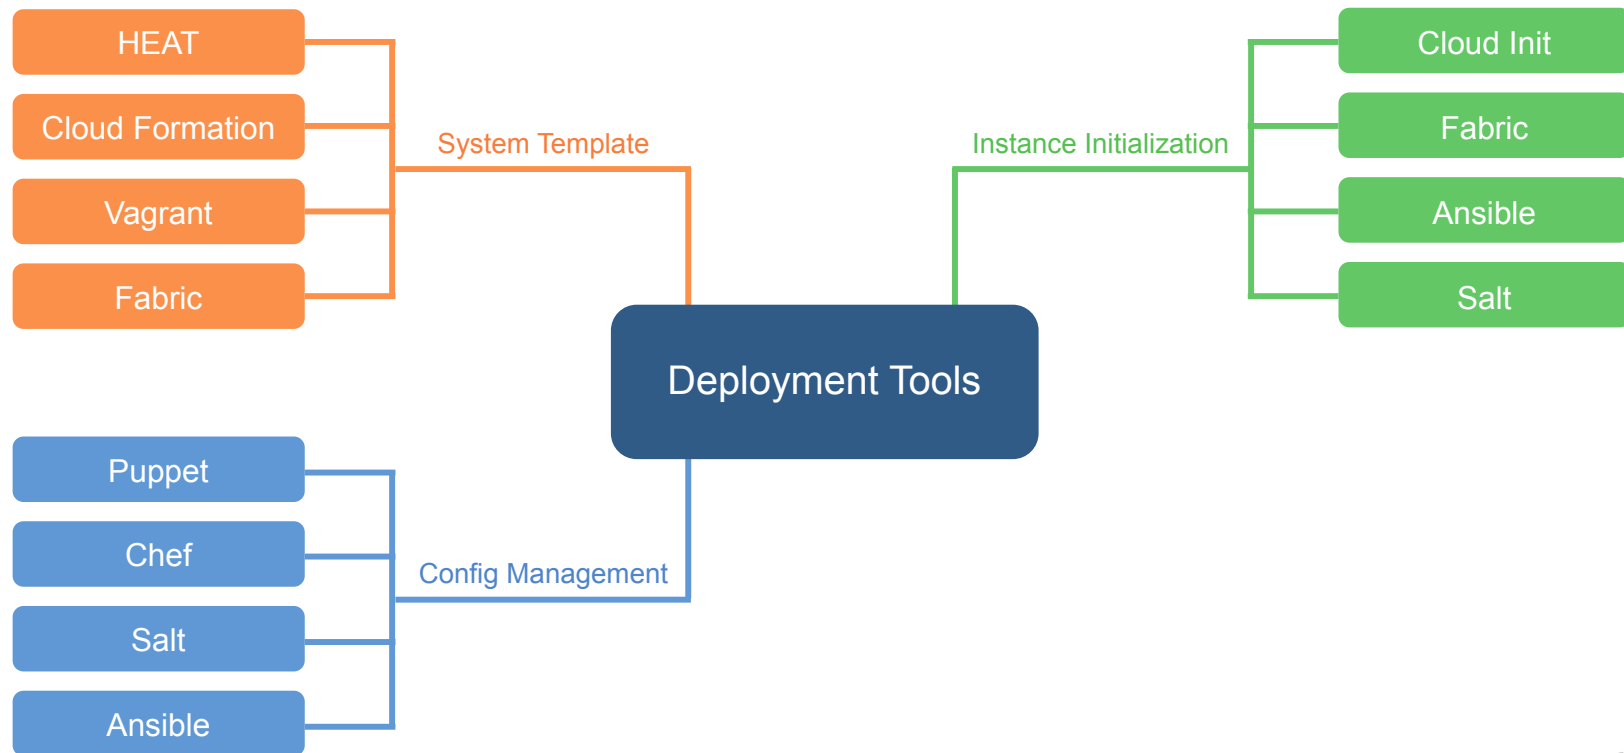
## Then

- Request VM (wait for IT to deploy)
- Load ISO, deploy OS
- Log in, install pre-compiled packages
- Configure Instance
- Request new REAL added to SLB
- 
- Sometimes
- Create golden image with app pre-set
- Deploy new copy, log in, tweak configs, restart services
- Request new REAL added to SLB

## Now

- Run Script...
  - Script launches generic image
  - Script passes init script to image
  - Image boots, runs init
  - Init loads application from repository
  - Config automated via puppet/chef
  - Script requests SLB update
- 
- Sometimes
- Ceilometer monitor adds/removes instances (based on script)

# Tool Map



# Step 1: Initialize

- Get basic info onto the new Instance
- SSH keys (avoid passwords)
- Perhaps update/install packages/fix DNS
- Download code (possibly)
- Download config management tools/components

# Cloud-Init Overview

- <https://cloudinit.readthedocs.org/en/latest/>
- Requires software on the “golden image”
- Embedded into Nova boot process and Neutron network setup
  - Metadata web service interface
  - Config\_drive mini-iso based model
  - Can pass script, yaml config, multi-part MIME message with scripts, config, etc.
- Principal Functions:
  - SSH key management (user, server)
  - User management
  - Package and Repository Management
  - Arbitrary Scripts
  - File creation

# Sample UserData as YAML

```
#cloud-config
write_files:
  - path: /etc/resolv.conf
    permissions: 0644
    owner: root
    content: |
      nameserver 8.8.8.8

apt_update: true
packages:
  - python-pip

  - git

  runcmd:
    - mkdir /home/ubuntu/
      gitrepo
    - git clone "https://
      github.com/mjhea0/flaskr-
      tdd.git" /home/ubuntu/
      gitrepo
    - pip install -r /home/ubuntu/
      gitrepo/requirements.txt
    - sed -i 's/app.run()/
```

app.run(host="0.0.0.0",port=1234)/g' /home/ubuntu/gitrepo/app.py

- python /home/ubuntu/gitrepo/app.py &

# Puppet

## What is it?

- Puppet is a data driven, idempotent configuration management system that allows you to define the state of your infrastructure and/or application and enforce the end state of the environment.

## What can it do?

- Manage packages (versions, installation, etc.)
- Manage configuration of components (add vhost to apache config)
- Report on state of configuration/ installation
- Enable a 'workflow' model for deployment



# Puppet hierarchy

- Module
  - Manifest
    - Init.pp
      - Class
      - Types/defines
- Example: class to define user

```
class robert ($state_val='present') {  
  user {"robert":  
    ensure=>$state_val,  
  }  
}
```

- In order to do something though, we

have to instantiate the class, so we can add:

```
class {'robert': state_val=>'absent'}
```

- And then we can “execute this (assuming we put it in a file test.pp)

```
puppet apply -vd init.pp
```

- Or we could put the file in /etc/puppet/modules/robert/manifests/init.pp and:  

```
puppet apply -ve 'class {"robert": }'
```

# Masterless Puppet

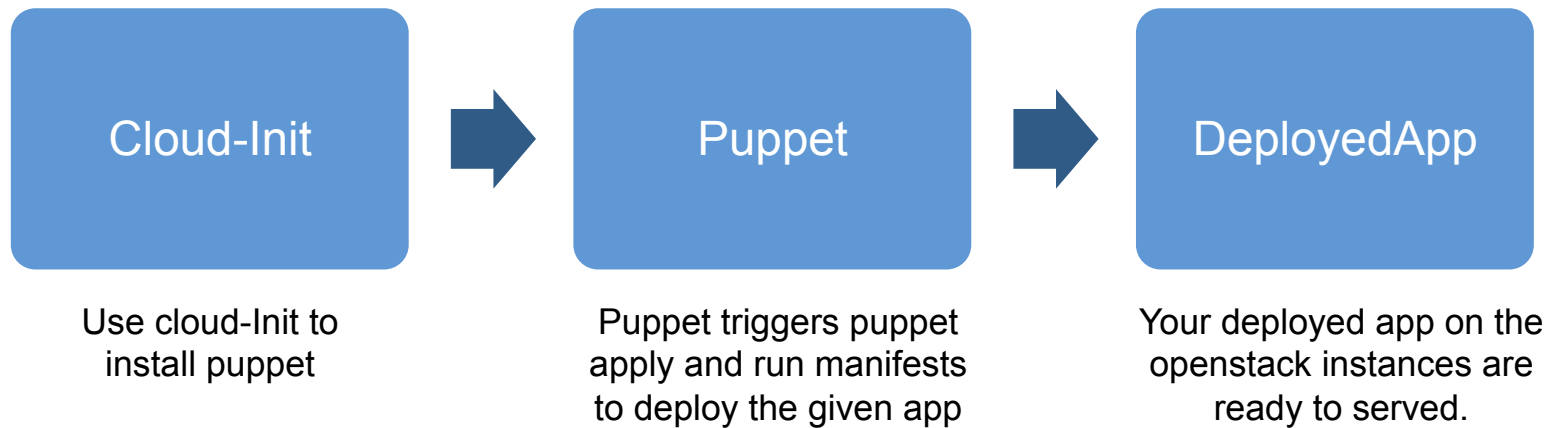
## **Advantages:**

- Masterless puppet modules are about as simple as writing bash scripts, however they are more maintainable, easily made to be cross platform.
- Still gives you the advantage of the abstraction layer that puppet provides

## **Disadvantages:**

- More difficult to distribute to nodes, need pulp, git or other distribution techniques

# How To Tie It All Together





# Using Heat and Ceilometer

# Ceilometer Overview

- Initial intent was to develop a tool to support billing against consumption
- Service Metering is a long standing solution, and has well understood models
- Data collected has uses besides metering
  - Alarm (something has happened)
  - Assurance (is my service alive)
  - Scale management (do I have too much/little)
  - Profiling (is my app running efficiently)



# Getting metrics from Ceilometer

- Principally via CLI, additional interactions are possible via HEAT ceilometer meters-list
- The list is long and includes more than you probably want to see, this is a more concise list:

Name	Type	Unit
cpu	cumulative	ns
cpu_util	gauge	%
disk.read.bytes.rate	gauge	B/s
disk.read.requests.rate	gauge	request/s
disk.write.bytes.rate	gauge	B/s
disk.write.requests.rate	gauge	request/s
network.incoming.bytes.rate	gauge	B/s
network.incoming.packets.rate	gauge	packet/s
network.outgoing.bytes.rate	gauge	B/s
network.outgoing.packets.rate	gauge	packet/s

# Ceilometer Alarm Notifier

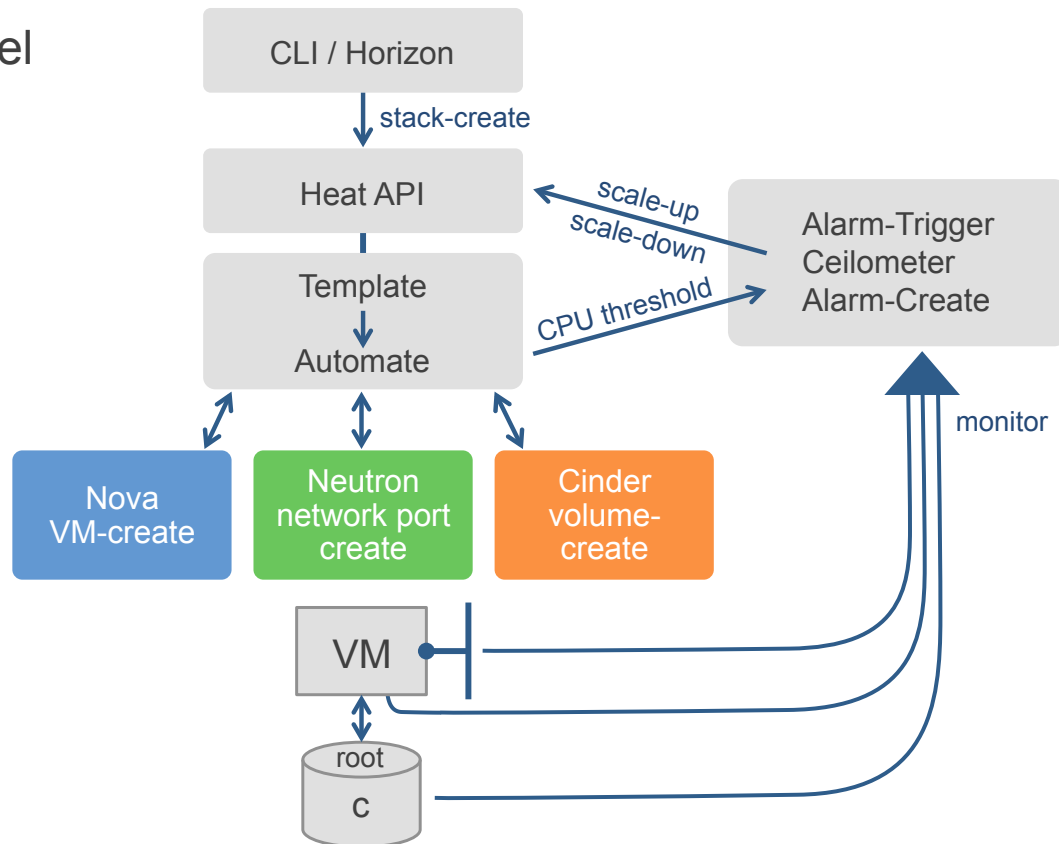
- Allows you to set alarms based on threshold evaluation for a collection of samples.
- An alarm can be set on a single meter, or on a combination
- To setup an alarm, you will call Ceilometer's API server specifying the alarm conditions and an action to take

```
ceilometer alarm-threshold-create --name cpu_high --description 'instance running hot' --meter-name cpu_util --threshold 70.0 --comparison-operator gt --statistic avg --period 600 --evaluation-periods 3 --alarm-action 'log://' --query resource_id=VM_ID
```

- Forms of action
  - HTTP/HTTPs Callback
  - Log

# Heat Model

## Heat Model





# Example HOT template

heat\_template\_version: 2013-05-23

description: Test Template

parameters:

ImageID:

type: string

description: Image use to boot a server

NetID:

type: string

description: Network ID for the server

resources:

server1:

type: OS::Nova::Server

properties:

name: "Test server"

image: { get\_param: ImageID }

flavor: "m1.tiny"

networks:

- port: { get\_resource: server1\_port }

server1\_port:

type: OS::Neutron::Port

properties:

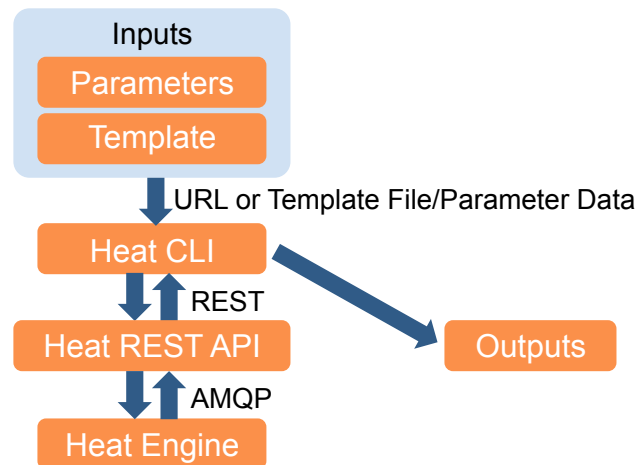
network\_ID: { get\_param: NetID }

outputs:

server1\_private\_ip:


description: IP address of the server in the private network

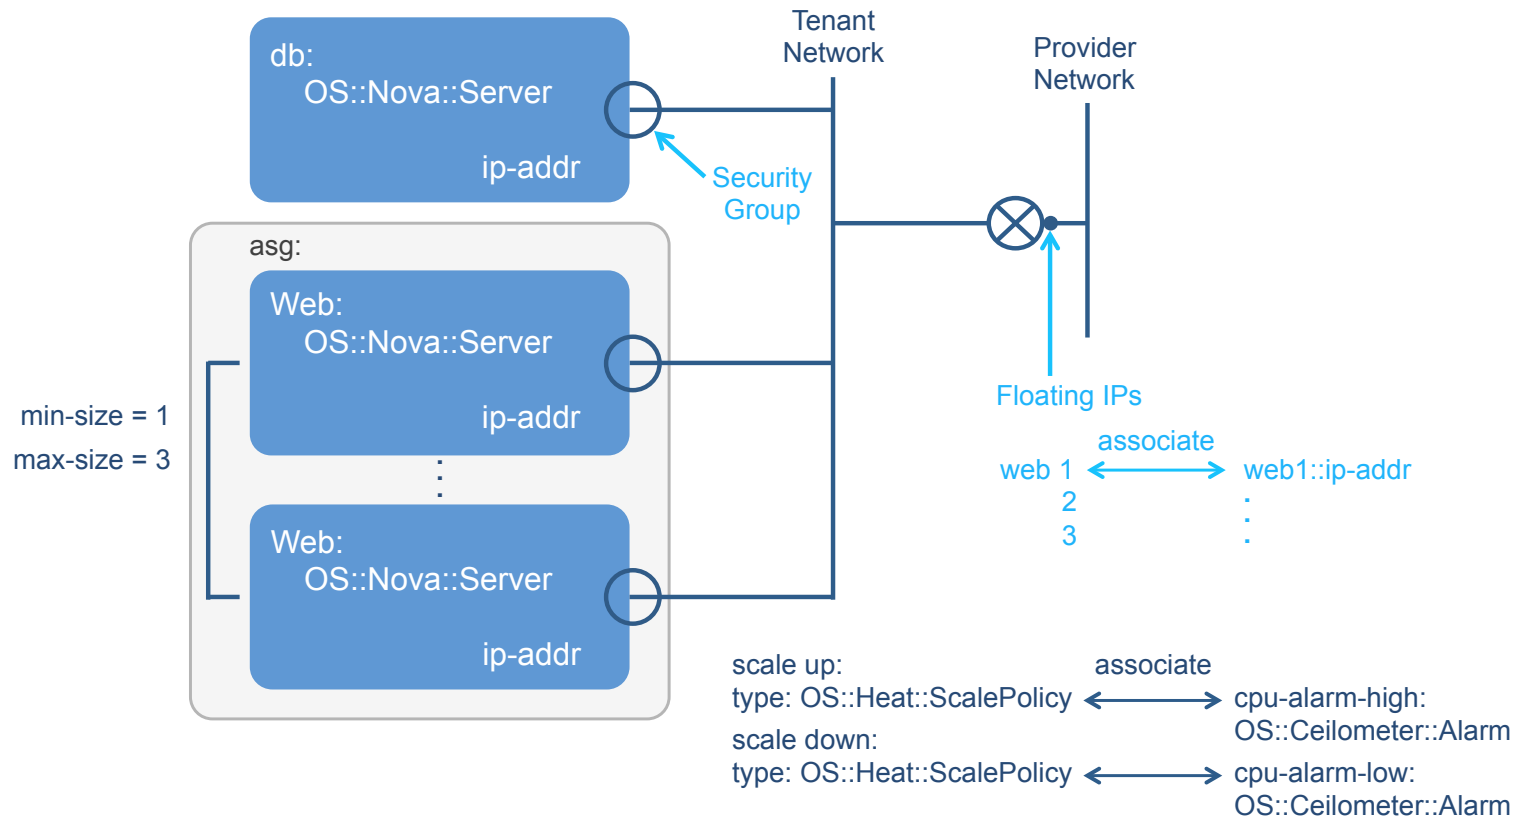
value: { get\_attr: [ server1, first\_address ] }



# Template Model considerations

- How your network works (floating Ips, SLB, etc.) still need to be considered
- Map the resources together to determine how pieces should fit
- Individual “public” VMs need Floating IPs associated
- Interaction between VMs (scale groups)
  - Scale up and Scale down resources
  - Associated Ceilometer alarms
- With SLB
  - Define pool (associate floating IP with pool)
  - Associate scale VMs with pool





# Example Scaling HOT

- From Github: <https://github.com/openstack/heat-templates/tree/master/hot>
- <https://raw.githubusercontent.com/openstack/heat-templates/master/hot/autoscaling.yaml>
- [https://raw.githubusercontent.com/openstack/heat-templates/master/hot/lb\\_server.yaml](https://raw.githubusercontent.com/openstack/heat-templates/master/hot/lb_server.yaml)

# Stack Creation

- `$ heat stack-list`

id	stack_name	stack_status	creation_time
4c712026-dcd5-4664-90b8-0915494c1332	stackone	CREATE_COMPLETE	2013-08-12T13:33:47Z
7edc7480-bda5-4e1c-9d5d-f567d3b6a050	stacktwo	CREATE_FAILED	2013-08-12T13:33:59Z

- `$ heat resource-list mystack`

logical_resource_id	resource_type	resource_status	updated_time
Database	OS::NOVA::Instance	CREATE_COMPLETE	2013-08-12T13:34:56Z

- `$ heat event-list mystack`

logical_resource_id	id	resource_status_reason	resource_status	event_time
Database	1	state changed	IN_PROGRESS	2013-08-12T13:42:19Z
Database	2	state changed	CREATE_COMPLETE	2013-08-12T13:45:44Z

# Network Topology

## Network Topology

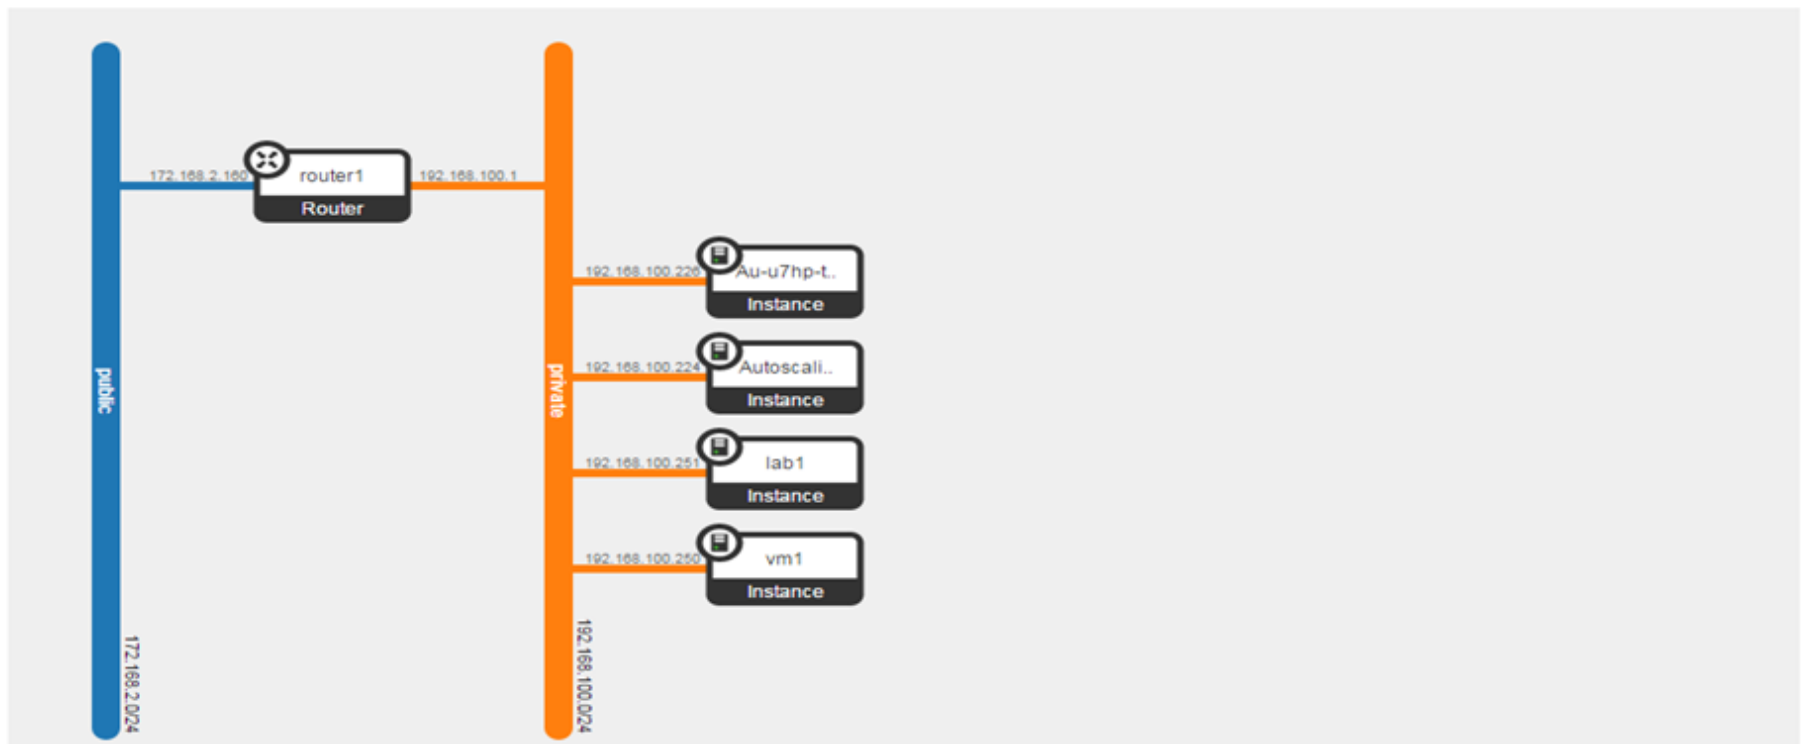
Small

Normal

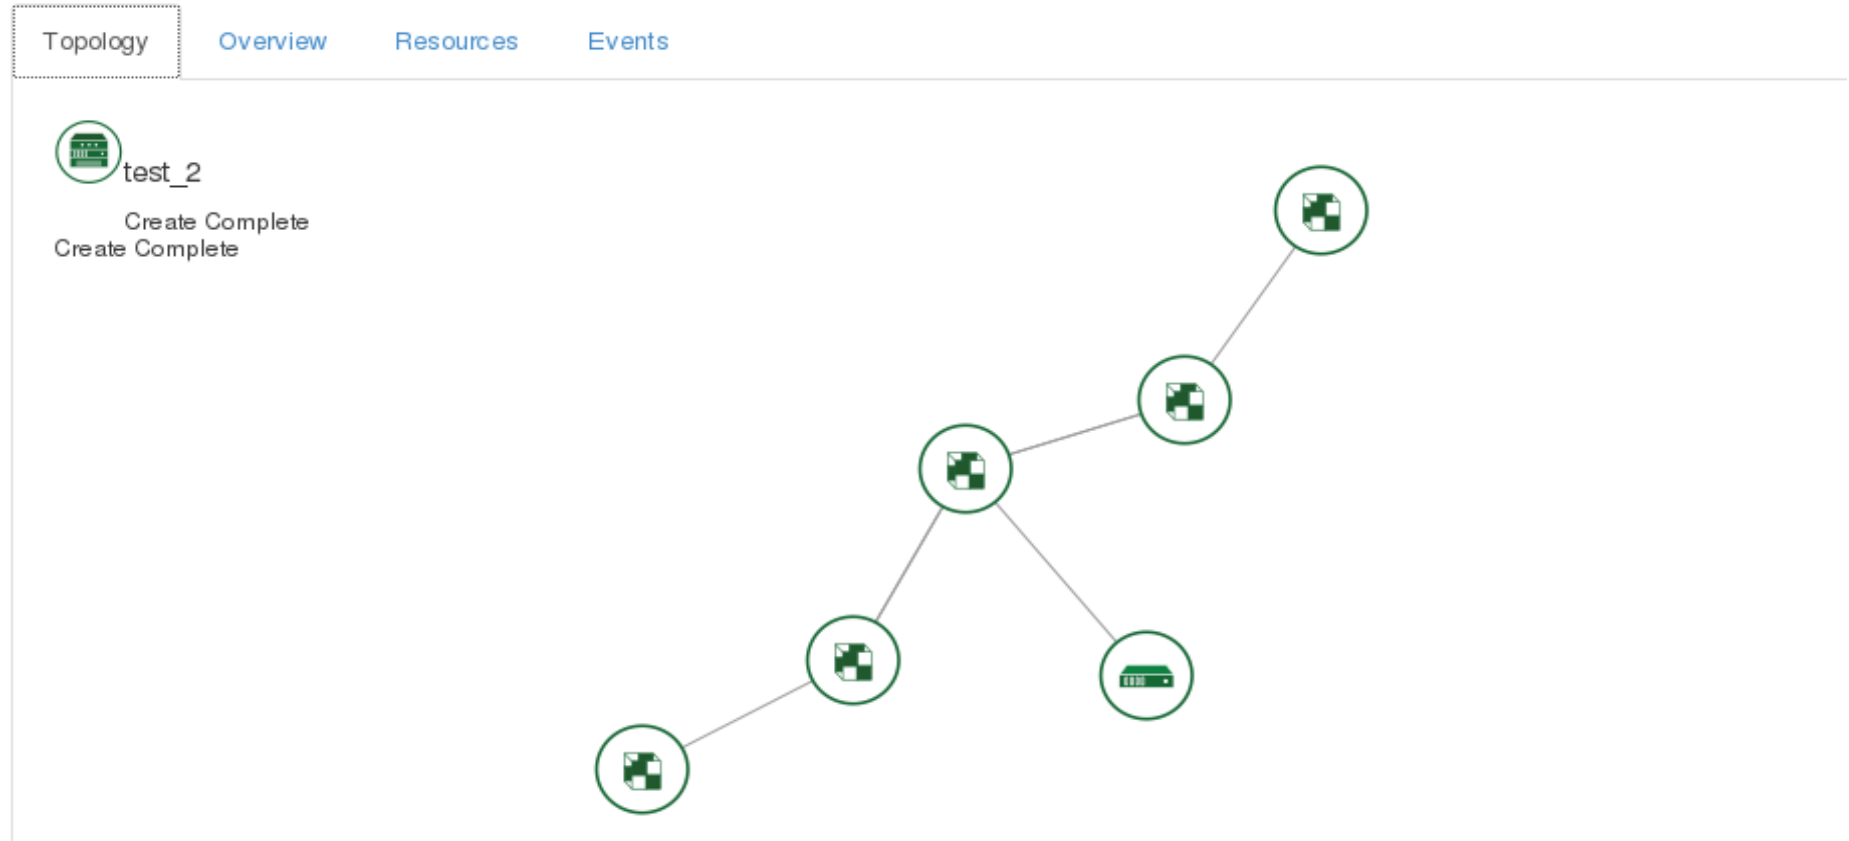
Launch Instance

Create Network

Create Router



# HEAT Topology



# Scaling without Ceilometer

- Can trigger HEAT scale up/scale down with URL from scale group.

- From example heat template:

scale\_up\_url:

value: { get\_attr: [web\_server\_scaleup\_policy, alarm\_url] }

scale\_dn\_url:

value: { get\_attr: [web\_server\_scaledown\_policy, alarm\_url] }

```
heat output-list stack
```

```
heat output-show stack scale_up_url
```

```
curl -X POST "http://10.1.10.13:8000/v1/signal/arn%3Aopenstack%3Aheat%3A%3Ae71d06beb59a40d1a9e29df6b01444e%3Astacks%2Fasg%2F8e386fdc-a872-493d-946b-c9f3e63878c3%2Fresources%2Fweb_server_scaleup_policy?Timestamp=2015-02-19T07%3A47%3A54Z&SignatureMethod=HmacSHA256&AWSAccessKeyId=bcd012555b7949bd953598f0090bb2cd&SignatureVersion=2&Signature=p3kbn%2F2v8zLiwUlsToobhPiHB4n8lRTfhnDsZMfbYhU%3D"
```



The logo features a green rectangular background. On the left side, there is a teal-colored cloud shape. The word "MONASCA" is written in white, bold, uppercase letters, with the "M" and "O" partially overlapping the teal cloud.

# MONASCA

## Monitoring as a Service - MONaaS

# Current Status

- Monasca and Stacktach.v3 is open-sourced in StackForge
- Not an OpenStack incubated project, but we are targeting incubation
- Metrics, Alarm Definitions, Alarms and Notification Methods completely supported/functional and ready for production deployment
- Who is working on it?
  - HP
  - RackSpace
  - IBM
- Who is deploying it?
  - HP: Public Cloud and Helion distribution
  - Time Warner Cable (TWC)
  - Workday

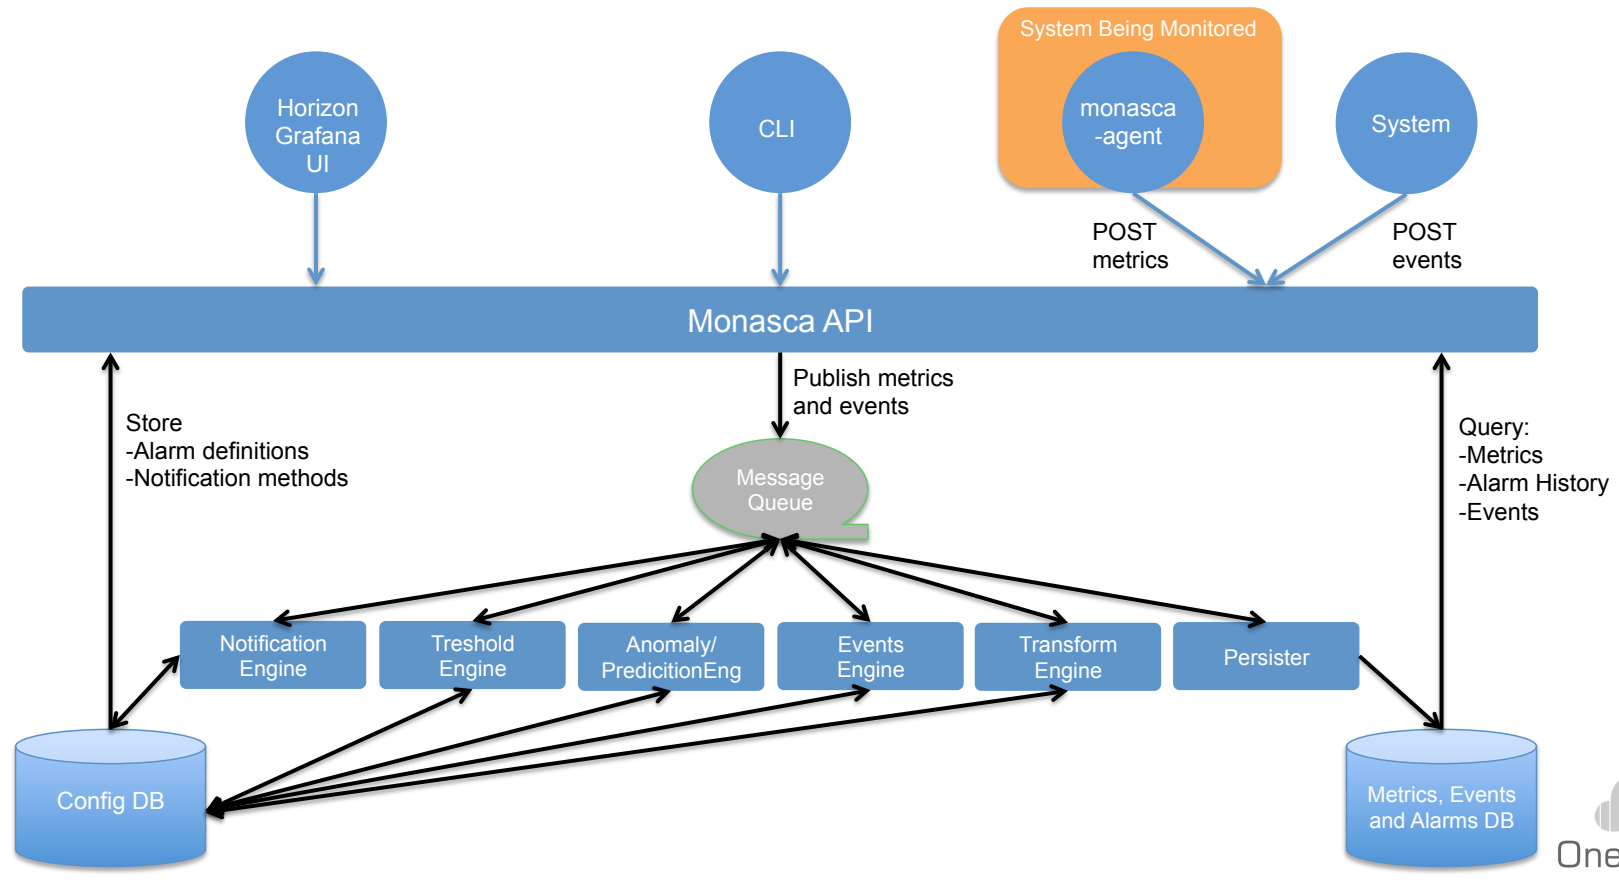
# Current Status- Progress

- Events/StackTach.v3 integration is in progress
- Anomaly detection is in progress
- Formalize micro-services architecture
  - Define message formats
  - Define how services are published and registered
- Python port is in progress:
  - All components Python except for API and Threshold Engine
  - API is 75% ported to Python. Note, Java API is 100% functional
  - Threshold Engine is the only remaining Java component

# What is Monasca?

- Monitoring-as-a-Service solution based on REST API
- Highly-performant, scalable, fault-tolerant and capable of big data retention
- Metrics storage/retrieval/statistics and alarm/thresholding engine
- Notification system
- Real-time event stream processing
- Open-source and built-on open-source technologies such as:
  - Kafka: Performant, scalable, fault-tolerant, durable message queue. Used by LinkedIn, Twitter, ...
  - Apache Storm: distributed realtime computation system
  - Time-series databases: InfluxDB supported today. Elastic-search in progress.
  - Consolidates multiple monitoring systems into a single solution
  - Used for both operational and customer facing monitoring.
- Extensible based on micro-services message bus architecture

# Architecture



# Metrics - Rest API

- Metrics: Create, query and get statistics for metrics

```
{  
  name: cpu.user_perc,  
  dimensions: {  
    hostname: hostname.domain.com,  
    region: uswest,  
    zone: 1,  
    service: compute  
  }  
}
```

Simple, concise flexible description  
Name (string)  
Dimensions: Dictionary of arbitrary (key, value) pairs

- Alarm Definitions –

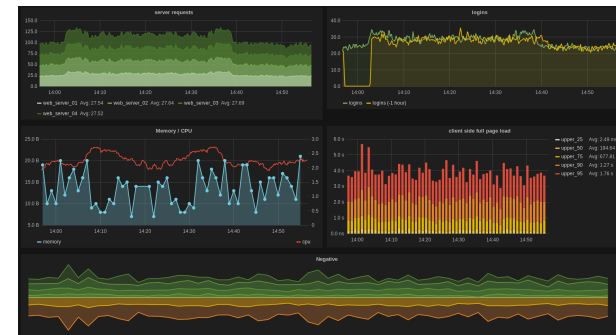
- Are templates that are used to automatically create alarms based on matching metric names and dimensions
- Simple compound expression grammar: `avg(cpu.user_perc{ }) > 85` or `avg(disk_read_ops{device=vda}, 120) > 1000`
- Actions associated with alarms for state transitions to ALARM, OK and UNDETERMINED
- Severity (LOW, MEDIUM, HIGH, CRITICAL).
- Alarms: Query and Delete alarms and query alarm state history
- Notification Methods: e.g. Email address. Associated with alarm definitions

# Metrics – monasca-agent

- Python monitoring agent
- System metrics (cpu, memory, ...)
- Service metrics (RabbitMQ, MySQL, Kafka, and many other)
- Application metrics (Built-in statsd daemon and Python Monasca Statsd library)
- VM metrics
- Active checks
  - HTTP status checks and response times
  - System up/down checks (ping and ssh)
- Runs any Nagios plugin
- Extensible/Pluggable: Additional services can be easily added

# Metrics - UI

- Horizon Dashboard
  - Overview/Top-level drill-down
  - Create/Read/Update/Delete alarm definitions using an expression builder
  - Read/Delete alarms and alarm history
  - Create/Read/Update/Delete notification methods
- Grafana Dashboard (<http://grafana.org/>)
  - Provides visualization of metrics

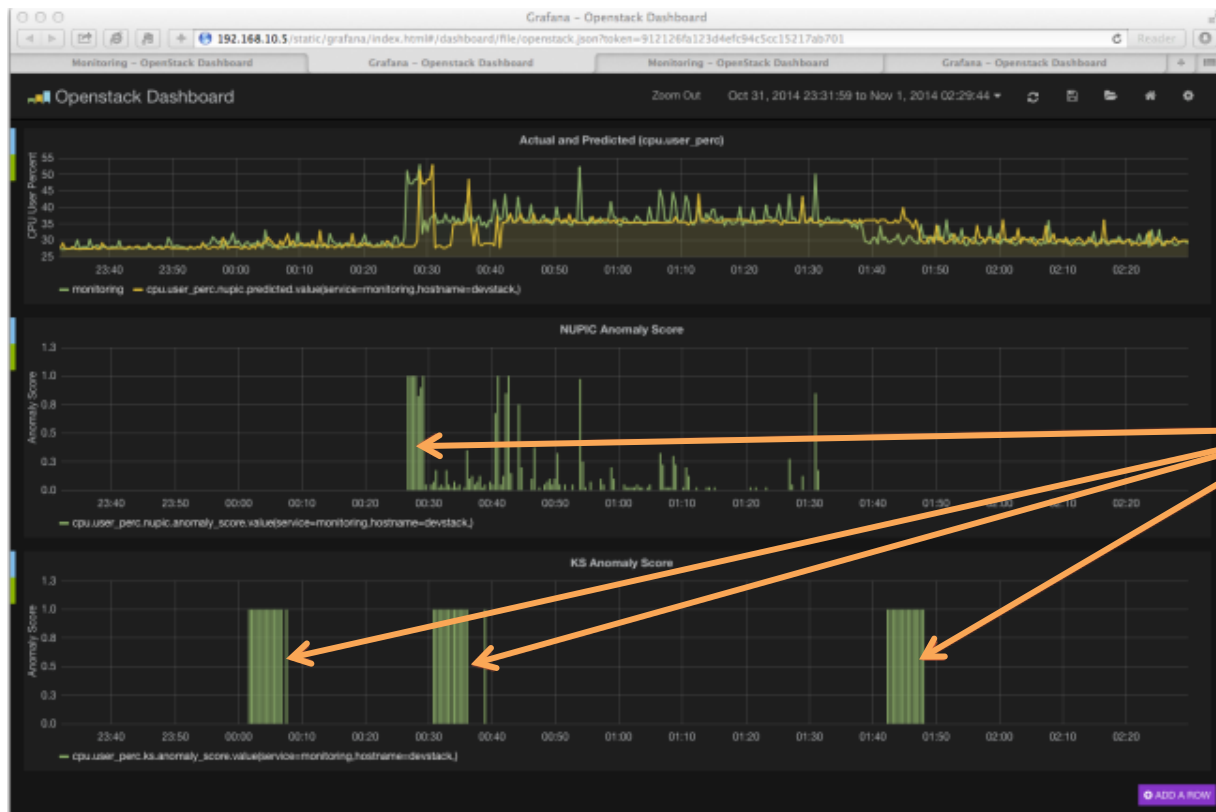




# Events – Anomaly Detection

- Monasca Anomaly Engine implements real-time streaming anomaly detection
- Two algorithms used:
  - Numenta Platform for Intelligent Computing (NuPIC) used by Grok  
(An open-source Python/C++ implementation of Hierarchical Temporal Memory)
  - Kolmogorov-Smirnov (K-S) Two Sample Test
- Anomaly Engine
  - Consumes metrics from the Kafka metrics topic
  - Calculates predicted value and anomaly score (probability of an anomaly)
  - Publishes calculated values as metrics to the Kafka metrics topic
- Alarms can be created for Anomaly scores

# Events – Anomaly Detection



Anomalies